

Advanced Application Development on Openshift

Prepared for: Mitzicom

Colum Foskin
cfoskin@redhat.com
November 2018

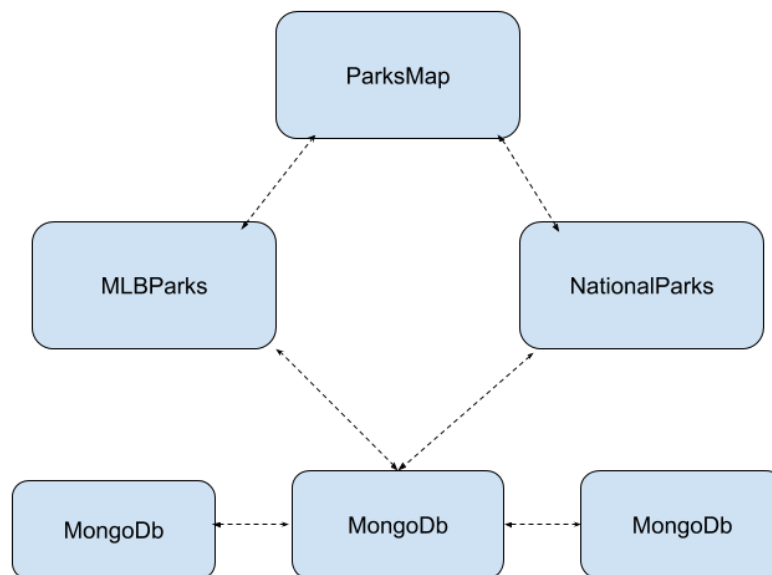
TABLE OF CONTENT

Introduction	2
OpenShift Projects Overview	3
Project Creation	3
Project Deployments	4
Continuous Integration & Continuous Deployment (CI/CD) Projects - Infrastructure	4
Jenkins	4
Nexus	5
SonarQube	5
Production & Development Projects	6
Development Project	6
Production Project	8
CI/CD Pipelines	9
MLBParks Pipeline	10
NationalParks Pipeline	10
ParksMap Pipeline	10
Project Source Code	11

Introduction

This is a Proof of Concept (POC) project using Red Hat OpenShift Container Platform (OCP) to determine the feasibility of using OpenShift as a target for an existing Java-based microservices workload. This project will demonstrate a fully integrated CI/CD pipeline using Nexus as the artifact repository and container registry and SonarQube for source code analysis. The microservices will be deployed to production in a blue-green strategy orchestrated by Jenkins. The application will consist of three microservices that will use a MongoDB replica set for storage, as shown in the following diagram.

Production Architecture



OpenShift Projects Overview

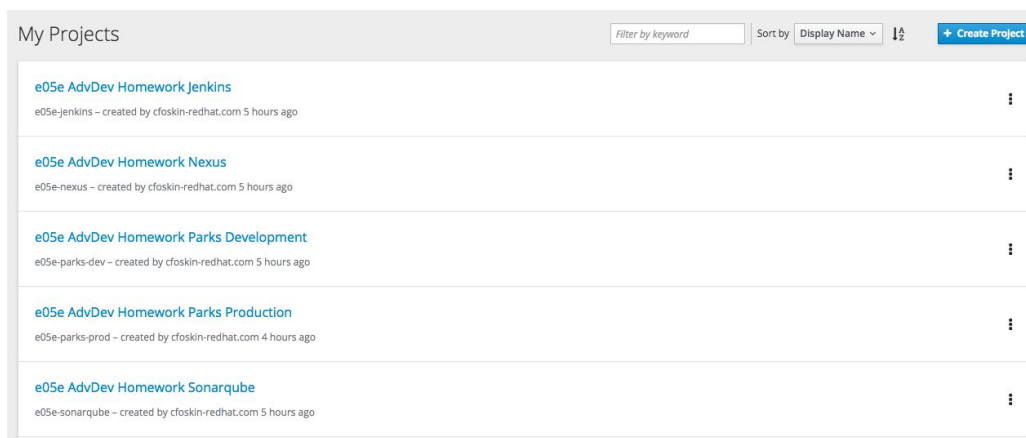
This POC will consist of five Projects deployed on OpenShift. The following section will detail what these will be and what purpose they serve. The projects will consist of most of OpenShifts key components or Objects such as:

- Pods
- Services
- Routes
- Deployments/DeploymentConfig
- ConfigMaps
- Secrets
- StatefulSets
- ImageStreams
- Pipelines
- Builds
- PersistentVolumeClaims

Project Creation

A script called ***create_projects.sh*** is provided to create the required projects. The following parameters are required when running this script: **GUID** , **USER**.

```
./setup_projects.sh GUID USER
```



Project Deployments

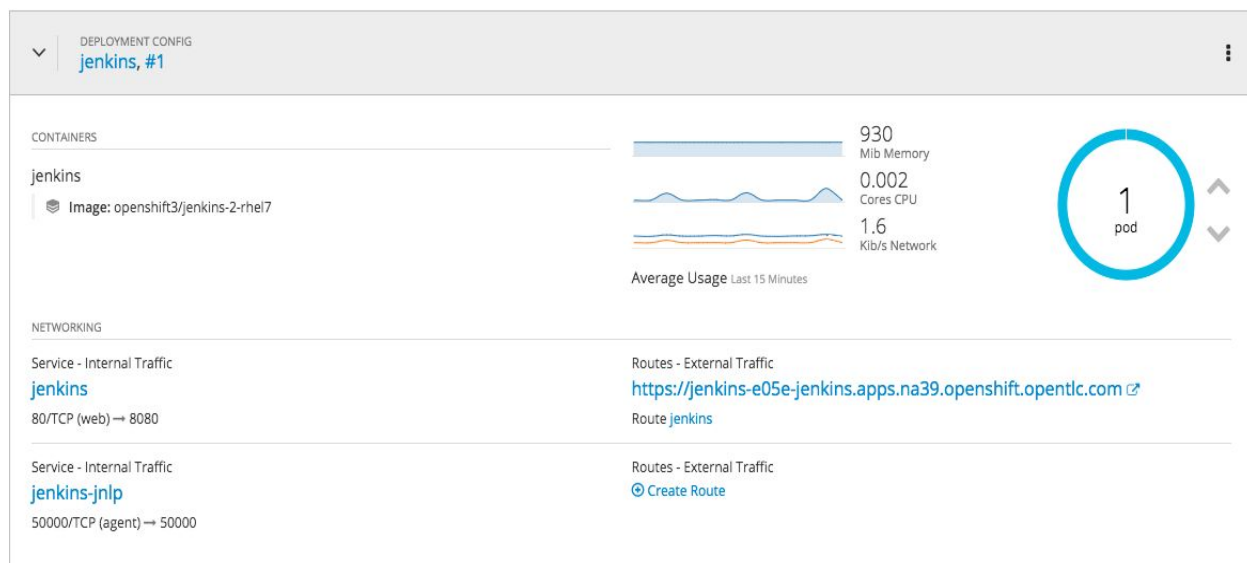
Continuous Integration & Continuous Deployment (CI/CD) Projects - Infrastructure

Jenkins

The Jenkins project consists of a Jenkins CI server to manage builds for each microservice. The Jenkins server uses a Persistent Volume for persistent storage. The Jenkins master will orchestrate builds to happen on a slave pod and will not perform the builds itself. The slave pods contain all of the required software to perform the Jenkins CI/CD pipeline stages.

A script called **setup_jenkins.sh** is provided to setup the jenkins project. The following parameters are required when running this script: **GUID , REPO & CLUSTER**.

```
./setup_jenkins.sh GUID REPO CLUSTER
```

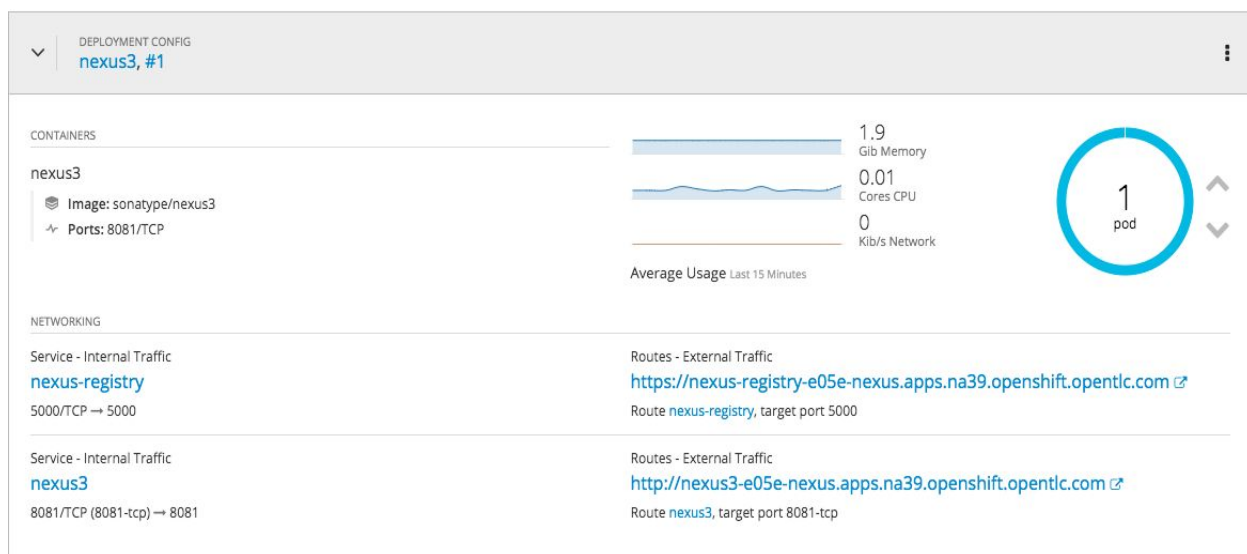


Nexus

The Nexus project will consist of a Nexus 3 server which will be used to cache Red Hat and other build artifacts and it will also be used as a Docker registry. Nexus 3 will be configured to use a Persistent Volume for persistent storage.

A script called **setup_nexus.sh** is provided to setup the Nexus project. The following parameters are required when running this script: **GUID**.

```
./setup_nexus.sh GUID
```



SonarQube

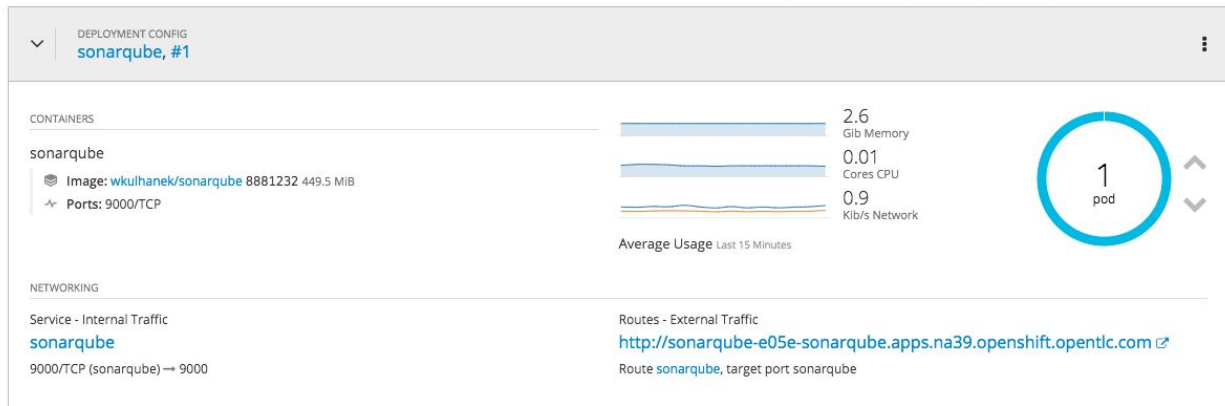
SonarQube will be used for analysing the code of each microservice and this will be deployed with a Postgresql instance which will be backed using a Persistent Volume for persistent storage.

A script called **setup_sonar.sh** is provided to create the SonarQube project. The following parameters are required when running this script: **GUID**.

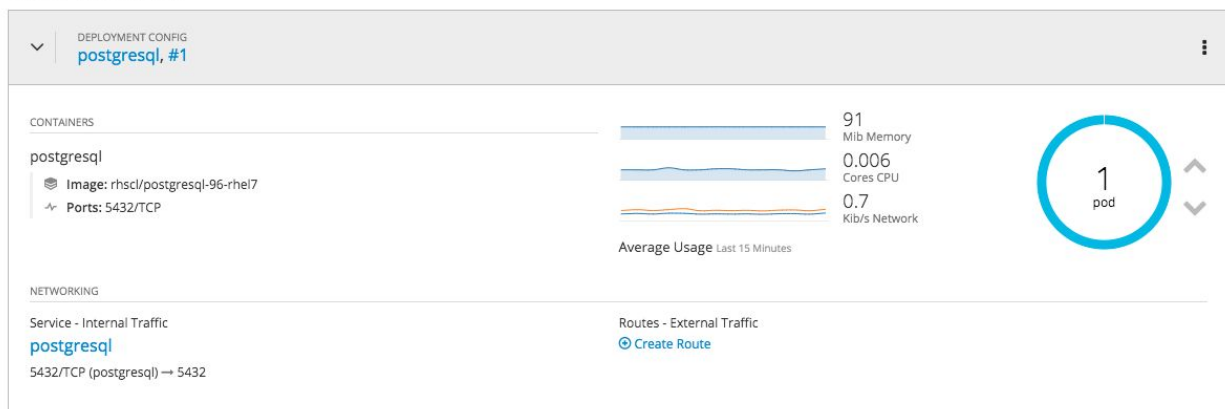
```
./setup_sonar.sh GUID
```

APPLICATION

sonarqube

<http://sonarqube-e05e-sonarqube.apps.na39.openshift.opentlc.com>

Other Resources



Production & Development Projects

Development Project

The development project will be used as a test environment for testing the microservices to conserve resources - instead of a separate “Test” project. This would also be the project that the developers would use for new features on the applications. This project will have a single persistent MongoDB pod for the microservices application.

A script called **setup_dev.sh** is provided to setup the Development project. The following parameters are required when running this script: **GUID**.

```
./setup_dev.sh e05e
```

APPLICATION

mlbparks

>

DEPLOYMENT CONFIG
mlbparks, #2

720
Mib Memory

< 0.01
Cores CPU

< 0.01
Kib/s Network

1 pod

⋮

APPLICATION

nationalparks

>

DEPLOYMENT CONFIG
nationalparks, #1

--
Mib Memory

--
Cores CPU

--
Kib/s Network

1 pod

⋮

APPLICATION

parksmapi

>

DEPLOYMENT CONFIG
parksmapi, #1

--
Mib Memory

--
Cores CPU

--
Kib/s Network

1 pod

⋮

Other Resources

>

DEPLOYMENT CONFIG
mongodb, #1

270
Mib Memory

0.01
Cores CPU

0.04
Kib/s Network

1 pod

⋮

Each microservice will be configured using a ConfigMap.

Name	Created	Labels
mlbparks-config	4 hours ago	none
nationalparks-config	4 hours ago	none
parksmapi-config	4 hours ago	none

Production Project

This project will consist of both blue & green versions of the microservices which will allow for blue-green deployment strategy. The microservices that run here will have come through the CI/CD pipelines and deemed to be ready for production.

A script called **setup_prod.sh** is provided to setup the Production project. The following parameters are required when running this script: **GUID**.

```
./setup_prod.sh GUID
```

APPLICATION

mlbparks-green

>

DEPLOYMENT CONFIG
mlbparks-green, #1

720 Mib Memory
< 0.01 Cores CPU
< 0.01 Kib/s Network

1 pod

⋮

APPLICATION

nationalparks-blue

>

DEPLOYMENT CONFIG
nationalparks-blue

No deployments for nationalparks-blue

⋮

APPLICATION

nationalparks-green

<http://nationalparks-e05e-parks-prod.apps.na39.openshift.opentlc.com>

>

DEPLOYMENT CONFIG
nationalparks-green, #1

--- Mib Memory
--- Cores CPU
--- Kib/s Network

1 pod

⋮

APPLICATION

parksmmap-blue

>

DEPLOYMENT CONFIG
parksmmap-blue

No deployments for parksmmap-blue

⋮

APPLICATION

parksmmap-green

<http://parksmmap-e05e-parks-prod.apps.na39.openshift.opentlc.com>

>

DEPLOYMENT CONFIG
parksmmap-green, #1

--- Mib Memory
--- Cores CPU
--- Kib/s Network

1 pod

⋮

Other Resources

>

STATEFUL SET
mongodb

390 Mib Memory
< 0.01 Cores CPU
1.8 Kib/s Network

3 pods

⋮

Each of the microservices are configured using ConfigMaps.

Name	Created	Labels
mlbparks-config-blue	4 hours ago	<i>none</i>
mlbparks-config-green	4 hours ago	<i>none</i>
nationalparks-config-blue	4 hours ago	<i>none</i>
nationalparks-config-green	4 hours ago	<i>none</i>
parksmapi-config-blue	4 hours ago	<i>none</i>
parksmapi-config-green	4 hours ago	<i>none</i>

CI/CD Pipelines

Each microservice application will have its own automated CI/CD Jenkins pipeline. The CI/CD pipeline will consist of the following stages:

- Check out the source code of the project.
- Build the war/jar file using Maven.
- Run unit tests using Maven.
- Perform a code analysis using SonarQube.
- Build the Openshift image & tag it.
- Publish the Maven artifact to the Nexus repository.
- Deploy the application to the Openshift Development project.
- Run a set of integration tests to ensure successful deployment.
- Copy the built image to the Nexus docker registry.
- Deploy the application to the Openshift Production project using the blue-green deployment strategy.
- Complete the switch over to the new version of the application after confirmation from the user.

The Pipelines are defined in a Jenkinsfile that is present in each microservice sub-directory of the application project. They are automatically created in Jenkins by using an Openshift BuildConfig with the JenkinsPipeline build strategy when the **setup_jenkins.sh** is ran to configure the Jenkins project.

The three pipelines and each stage that is executed can be seen in the following screenshots.

MLBParks Pipeline

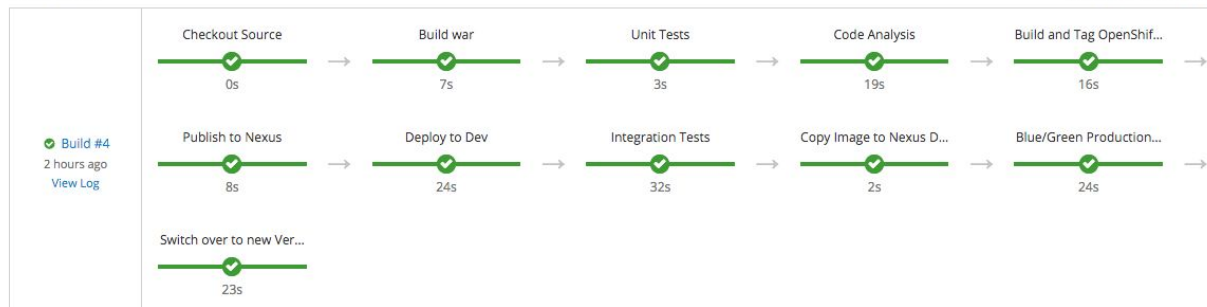
mlbparks-pipeline created 3 hours ago

[Start Pipeline](#)

Source Repository: https://github.com/cfoskin/appvdev_homework.git

Recent Runs

Average Duration: 4m 56s



[View Pipeline Runs](#) | [Edit Pipeline](#)

NationalParks Pipeline

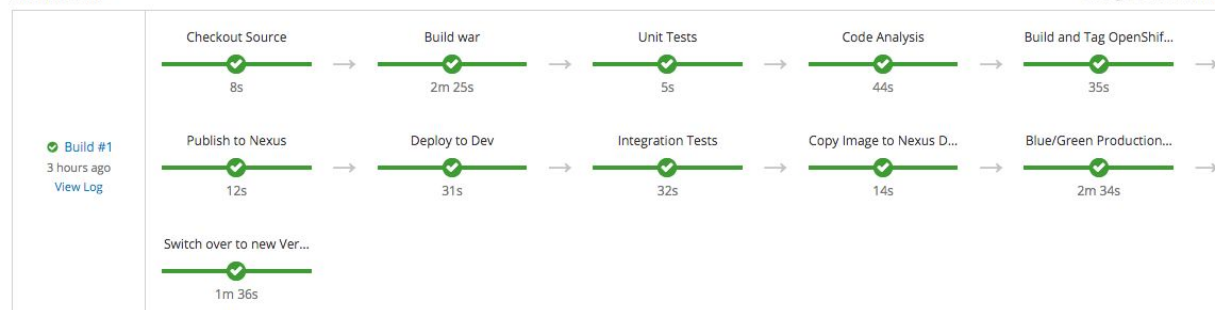
nationalparks-pipeline created 3 hours ago

[Start Pipeline](#)

Source Repository: https://github.com/cfoskin/appvdev_homework.git

Recent Runs

Average Duration: 12m 1s



[View Pipeline Runs](#) | [Edit Pipeline](#)

ParksMap Pipeline

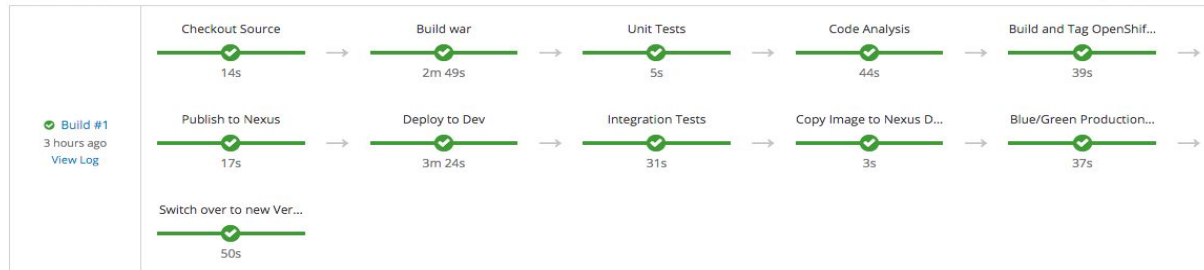
parksmapi-pipeline created 3 hours ago

[Start Pipeline](#)

Source Repository: https://github.com/cfoskin/appvdev_homework.git

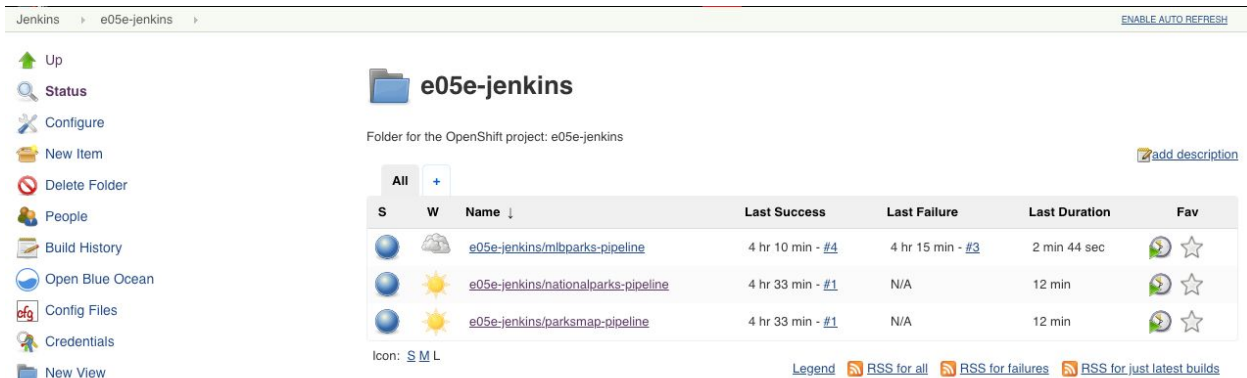
Recent Runs

Average Duration: 12m 8s















[View Pipeline Runs](#) | [Edit Pipeline](#)

The pipelines can also be viewed or managed via the Jenkins console.



Folder for the OpenShift project: e05e-jenkins

S	W	Name ↓	Last Success	Last Failure	Last Duration	Fav
		e05e-jenkins/mlbparks-pipeline	4 hr 10 min - #4	4 hr 15 min - #3	2 min 44 sec	 
		e05e-jenkins/nationalparks-pipeline	4 hr 33 min - #1	N/A	12 min	 
		e05e-jenkins/parksmapi-pipeline	4 hr 33 min - #1	N/A	12 min	 

Icon: [S](#) [M](#) [L](#)

[Legend](#) [RSS for all](#) [RSS for failures](#) [RSS for just latest builds](#)

Project Source Code

The source code for this project can be found in a GitHub repository [1]. This includes all of the scripts and Openshift templates needed to setup and configure each project. It also includes the code for each microservice where each one has its own Jenkinsfile with the pipeline definition.

[1] https://github.com/cfoskin/appvdev_homework